

Scheduling of Serverless Functions

Latest researches and approaches

By: Aryan Ebrahimpour

Professor: Dr. Entezari Maleki



Table of Contents

Introduction

-  What are FaaS and Serverless?

-  Challenges

-  Goals

-  Platforms

Approaches

-  Categories

-  Cold-start approaches

-  Scientific workflow approaches

-  Parallel approaches

References

What are FaaS and Serverless?

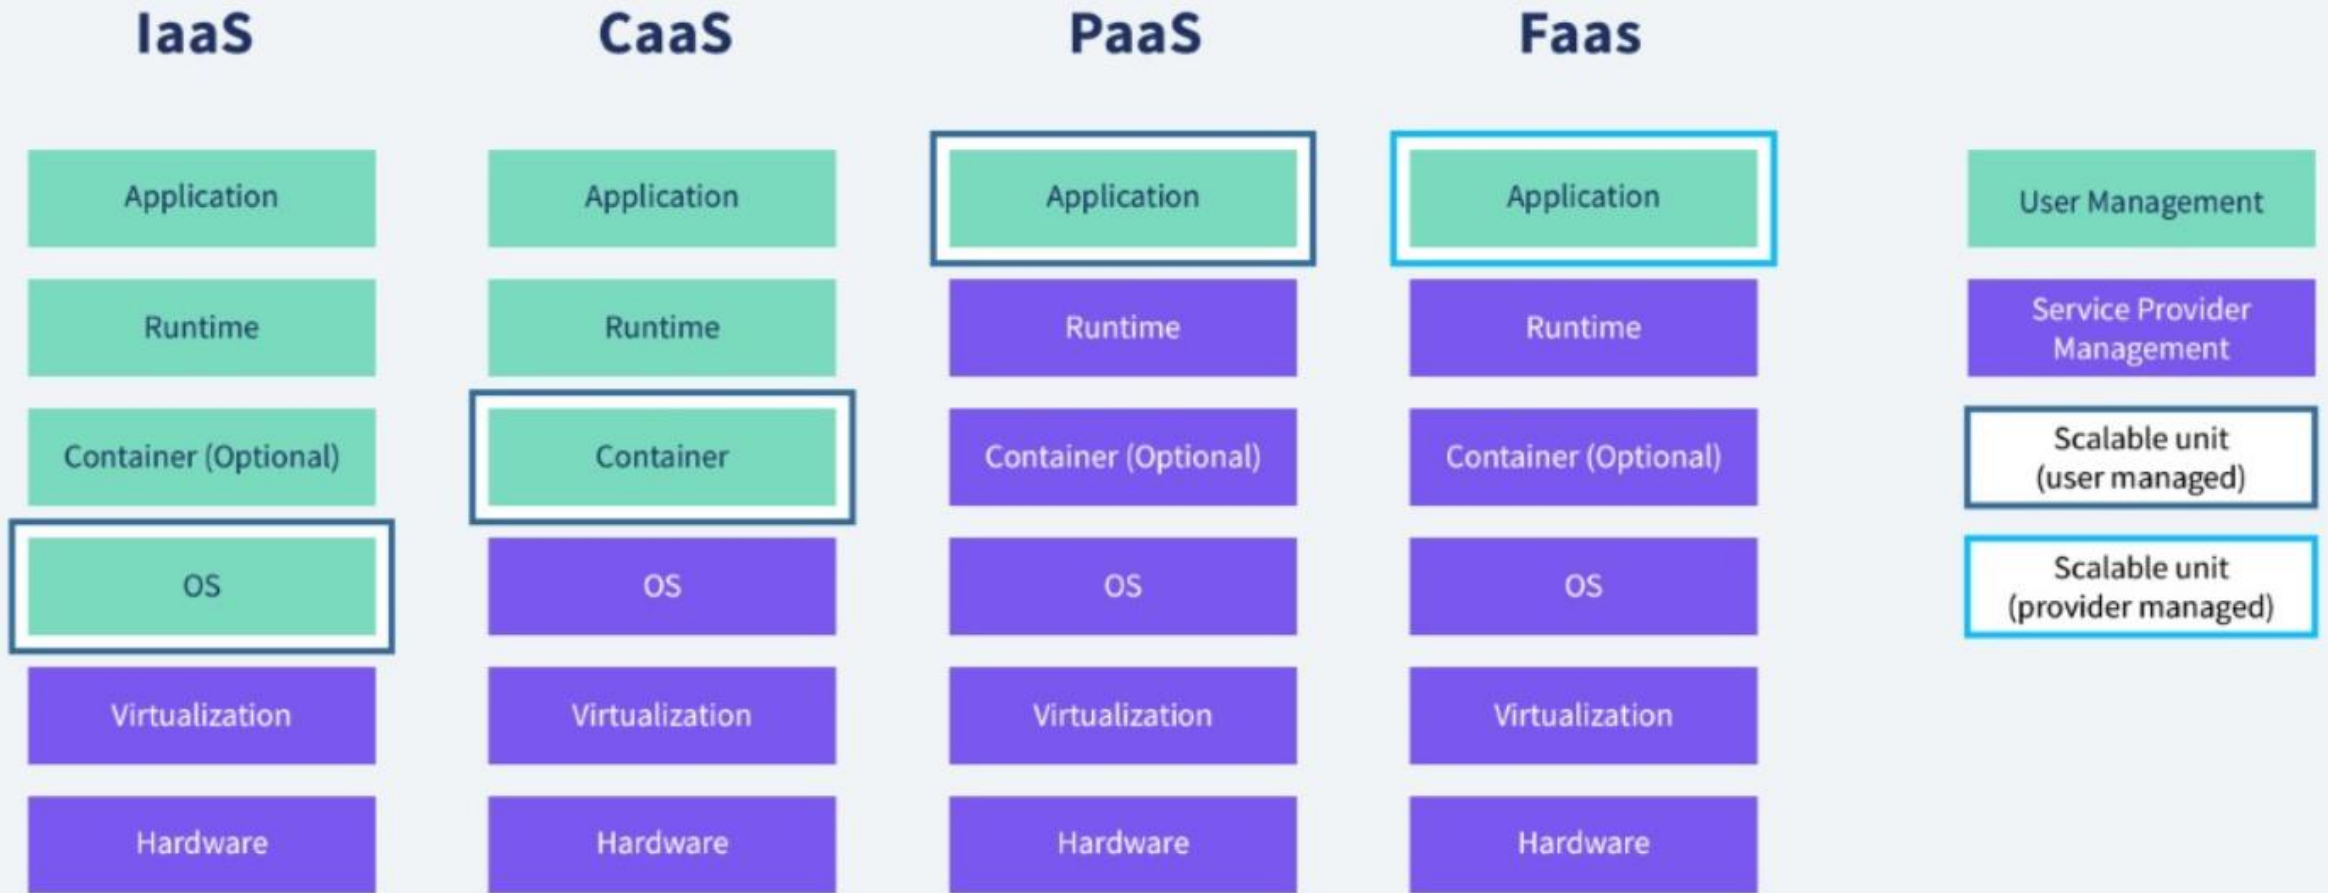
Provider-User scopes, and FaaS vs BaaS



What are FaaS and Serverless?

- 📖 FaaS stands for “Function as a Service”.
- 📖 Its popularity comes from popularity of Containers and Microservices.
- 📖 Industrialized by Amazon with Amazon Lambda back in 2014.

What are FaaS and Serverless?



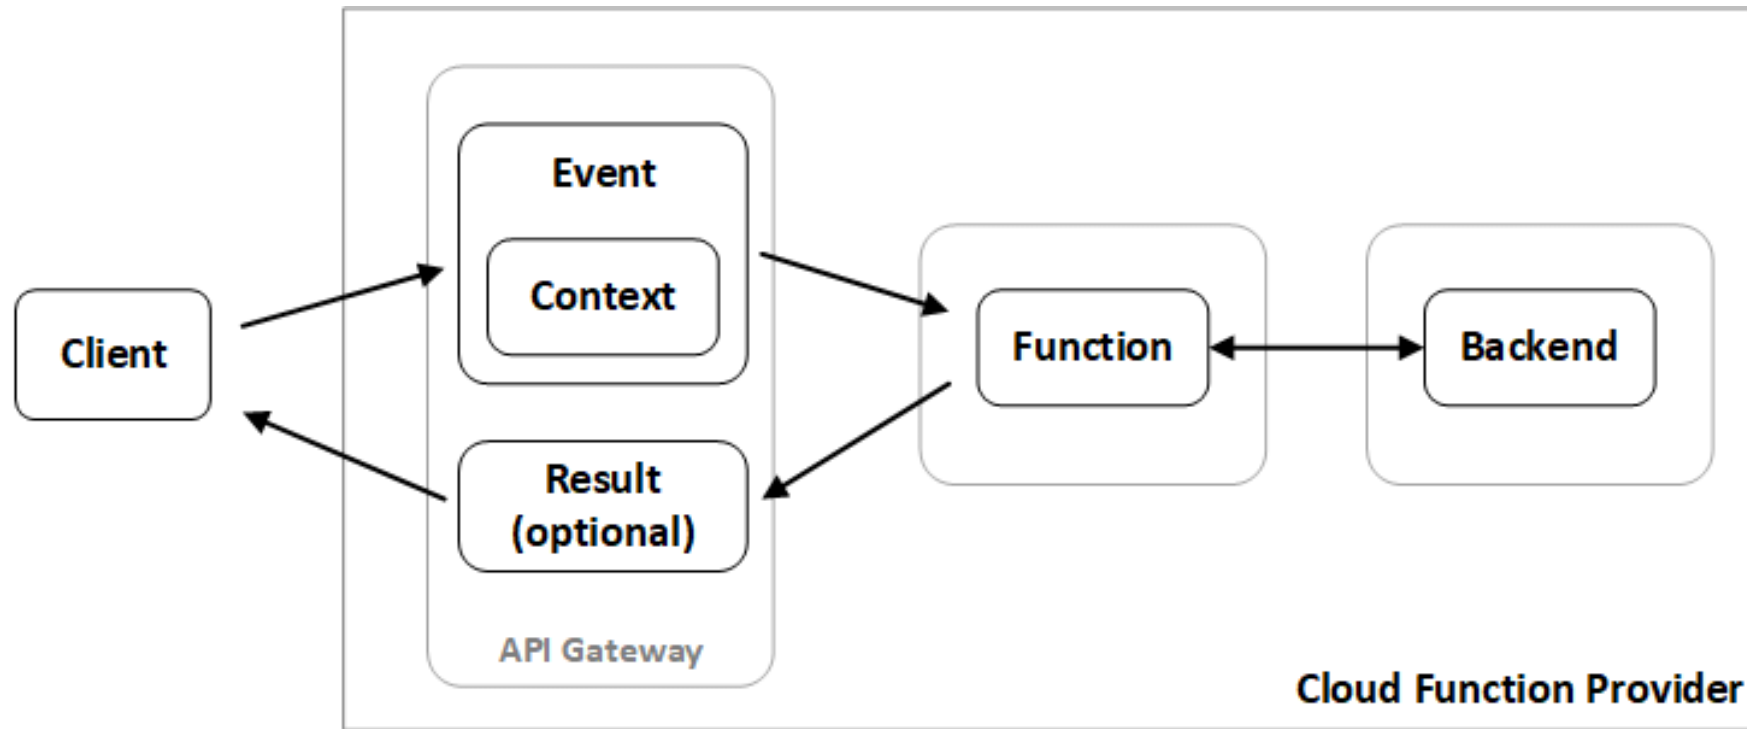
fine-grained functions instead of coarse-grained apps

What are FaaS and Serverless?

- 📖 The terms “**Serverless**” and “**FaaS**” are used interchangeably.
- 📖 However they are not quite the same.
- 📖 Serverless doesn’t mean you don’t have servers.
- 📖 FaaS is one type of Serverless computing.

What are FaaS and Serverless?

📁 FaaS is **event-driven**.



What are FaaS and Serverless?

- 📖 Don't confuse it with "BaaS".
- 📖 In BaaS, you'll develop your front app, while having the backend already up and running in the cloud.
- 📖 Some say BaaS and FaaS are both serverless computing.

What are FaaS and Serverless?

Frontend

(Developer builds)

- User interface
- Client-side logic



Backend

(Vendor provides as a service)

- Database management
- Cloud storage
- User authentication
- Push notifications
- Hosting

Firestore, is that you?

Challenges

Why scheduling for FaaS matters?



Challenges

- 📖 You might say we already have scheduling approaches that actually work.
- 📖 Then why do we need FaaS specific scheduling systems?
- 📖 Other methods don't seem to very well fit this model.

Challenges

- 📖 Cloud platforms offer cheap and scalable services.
- 📖 However FaaS should be **even more scalable**, and **less costlier** than other services.
- 📖 This will end up with challenges that provider should address.

Challenges

Cost

As it should be cheaper, scheduler should be even more efficient by not running for excess time. Warm up time is considerable.

Challenges

Isolation

Functions are stateless. They must be isolated from environment even more than other models.

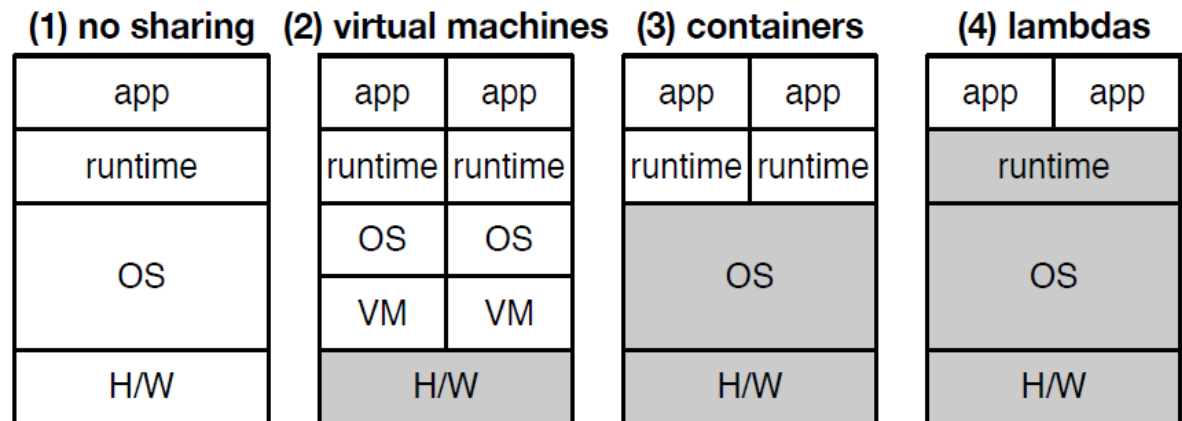
Orchestration and Load balancing

Prediction of function calls must be considered to have a better management on the number of up and running containers for running the function.

Challenges

Security

Many functions from different customers might be running on a single physical or virtual machines, and sometimes even on a container forked from the same base container. This leads to some security concerns.



Challenges

Scheduling

A good scheduling algorithm directly affects **Cost**, **Response Time** and **Scalability** of these cloud functions.

Challenges

Others

There are also other challenges like vendor lock-in problem, old software, IDEs and dev tools integration, state management, etc.

However we focus on Scheduling.

Goals

What are we aiming to improve with these algorithms?



Cold-start Time

This time includes:

1. Fetching information
2. Creating container
3. Startup time
4. Loading packages time

Response Time

Time spent since activating the function (usually from an HTTP event) until the final response.

Throughput

Number of responses received from the cloud function provider in a unit of time of a simulated environment.

Final cost

Even though it is highly affected by the other factors that are already mentioned, algorithm can also consider the user budget.

Platforms

Some of the platforms



Source	Site	Owner	Platform Name
proprietary	https://aws.amazon.com/lambda	Amazon	Amazon Lambda
proprietary	https://azure.microsoft.com/en-us/services/functions	Microsoft	Azure Functions
proprietary	https://cloud.google.com/functions	Google	GCP Functions
open source	https://kubernetes.io/docs/concepts/overview/kubernetes-api/	kubeless	kubeless
open source	https://www.openfaas.com	OpenFaaS Ltd	Open FaaS
proprietary	https://www.ibm.com/uk-en/cloud/functions	IBM	IBM Cloud Functions
open source	https://openwhisk.apache.org	Apache	Open Whisk
open source	https://fission.io	Platform9	Fission
open source	https://open.iron.io	Iron	IronFunctions
open source	https://fnproject.io	Oracle	Fn Project
open source	https://knix.io	Nokia Bell Labs	KNIX
open source	https://github.com/open-lambda	OpenLambda	Open-lambda
open source	https://knative.dev	Google	Knative
proprietary	https://www.alibabacloud.com/products/function-compute	Alibaba	AlibabaCloud Function Compute
proprietary	https://vercel.com/docs/serverless-functions/introduction	Vercel	Vercel Serverless Functions
open source	https://nuclio.io	Iguazio	Nuclio
proprietary	https://www.slappforge.com/sigma	SLAppForge Inc	Sigma


Why even bother knowing the platforms?

HOW STANDARDS PROLIFERATE:

(SEE: A/C CHARGERS, CHARACTER ENCODINGS, INSTANT MESSAGING, ETC.)

SITUATION:
THERE ARE
14 COMPETING
STANDARDS.

14?! RIDICULOUS!
WE NEED TO DEVELOP
ONE UNIVERSAL STANDARD
THAT COVERS EVERYONE'S
USE CASES.



SOON:

SITUATION:
THERE ARE
15 COMPETING
STANDARDS.

Approaches

In what ways people tried to address these challenges?



Approaches

📖 **Cold Start** approaches: address cold start via optimized package loading

📖 Optimization for **Scientific workflows**

📖 **Parallel** approaches

Cold-start approach

Approaches > Cold Start

- 📦 Loading packages and software dependencies take a lot of time for starting container.
- 📦 Technologies that are facing dependency-hell problems.


Approaches > Cold Start

- 📦 Harter et al., changed the **Linux kernel** to optimize containers performance.
- 📦 These container preload the packages that are needed to run the software, and **lazy load** the others.
- 📦 They've shown **76% of the startup time** is related to package loading.
- 📦 They've succeeded to increase **production cycle by 5x** and **development cycle 20x**.
- 📦 Many systems can use this to boost their startup time.

Approaches > Cold Start

- 📖 Akkus et al., proposed SAND. A model of isolation to address cold-start.
- 📖 They modeled the problem to Apps and Functions.
- 📖 Apps are in different containers while their functions are in the same container, but different processes by forking the warm container.
- 📖 Some other methods took this idea, and created a warm python interpreter with the most used libraries preloaded.


Approaches > Cold Start



Approaches > Cold Start

- 📦 Aumala et al., also addressed huge packages and libraries problem.
- 📦 They bundled the packages with compute nodes instead of functions.
- 📦 Functions run on the nodes that already have the most heavy package loaded in cache.
- 📦 However, this method only considers the heaviest package.

Approaches > Cold Start




 Aumala's method is based on two algorithms:

1. Consistent Hashing
2. Power of two choices technique

Approaches > Cold Start

consistent hashing,
scalable & easy to use



📖 Aumala's method is based on two algorithms:

1. Consistent Hashing

2. Power of two choices technique

📖 Power of two choices selects two random nodes if there are multiple candidates, then picks the least loaded one.

📖 If the selected machine is overloaded already, algorithm fallbacks to the least loaded mechanism.

Algorithm 1: Package-aware scheduler algorithm (PASch)

Global data: List of workers, $W = \{w_1, \dots, w_n\}$, and their load thresholds, $T = \{t_1, \dots, t_n\}$, mapping function M

Input: Function, f , largest required package, p

```
1 if ( $p$  is not null)then
    /* Get affinity workers */
2     $\langle a1, a2 \rangle = M(p)$ 
    /* Select target with least load */
3    if ( $load(w_{a1}) < load(w_{a2})$ )then
4         $A := a1$ 
5    else
6         $A := a2$ 
    /* If target is not overloaded, we are done */
7    if ( $load(w_A) < t_A$ )then
8        Assign  $f$  to  $w_A$ 
9        return
    /* Balance load */
10 Assign  $f$  to least loaded worker,  $w_i$ 
```

Algorithm 2: Mapping function, M ; given a package, returns two affinity nodes.

Global data: A consistent hash implementation, *consistent*, and value to be added to the package ID to map a second affinity worker to it, $salt$

Input: Package id, p

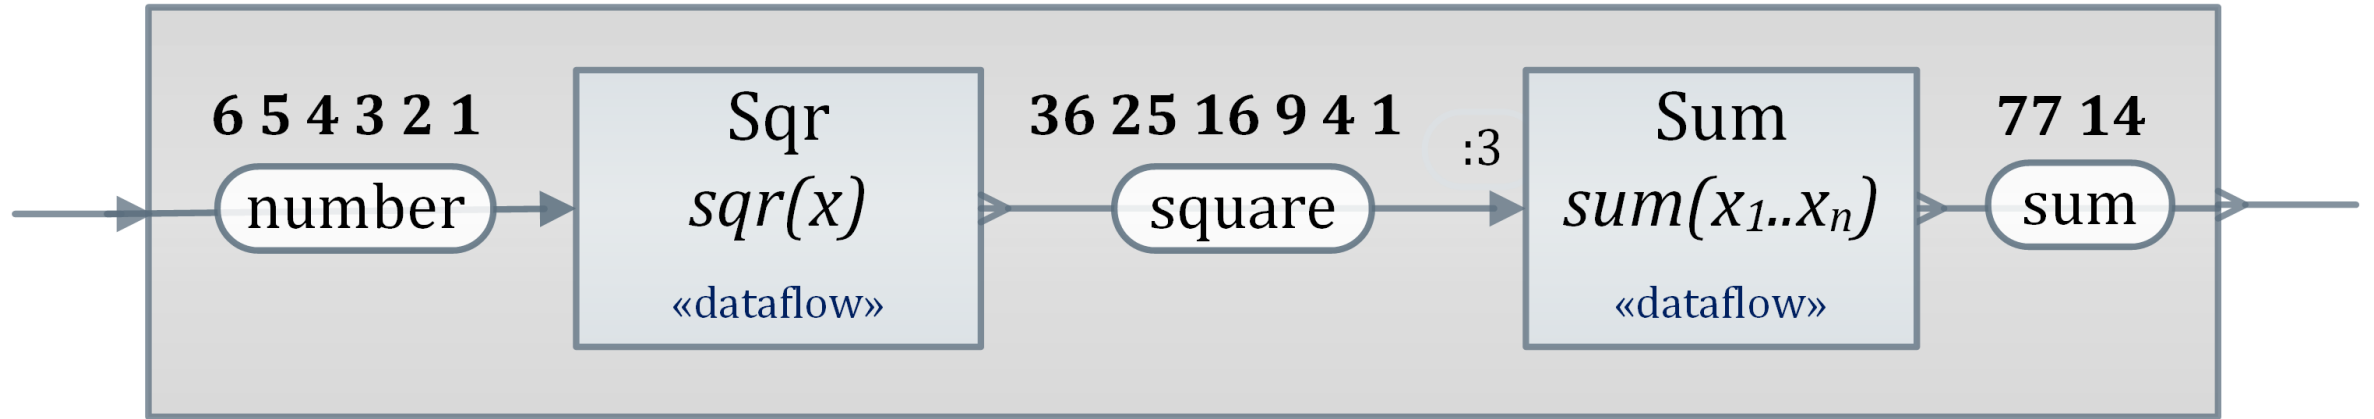
Output: Affinity workers for p , $\langle a1, a2 \rangle$

```
/* Get two affinity workers */
1  $a1 = \text{consistentHash.get}(p)$ 
2  $a2 = \text{consistentHash.get}(p + salt)$ 
3 return  $\langle a1, a2 \rangle$ 
```

Scientific Workflow approach

Approaches > Scientific Workflow

- 📖 As scientific workflows are important use case of FaaS, there are approaches to specifically address these type of cloud functions.
- 📖 But what are Scientific Workflows?
- 📖 They are a set of computations (usually mathematical) that are represented as **DAGs**.
- 📖 $G = \langle T, E, \text{Data} \rangle$ where $T = \{t_1, t_2, \dots\}$ are tasks, and E are dependency edges.
- 📖 The E set ensures that a child task shouldn't start before its parent finishes.




Approaches > Scientific Workflow

- 📖 Malawski et al., analyzed different providers to see if their cloud function services fit for scientific workflows.
- 📖 They developed a scientific workflow runner on them.
- 📖 They've shown 5 models and options for developing such systems.

Approaches > Scientific Workflow

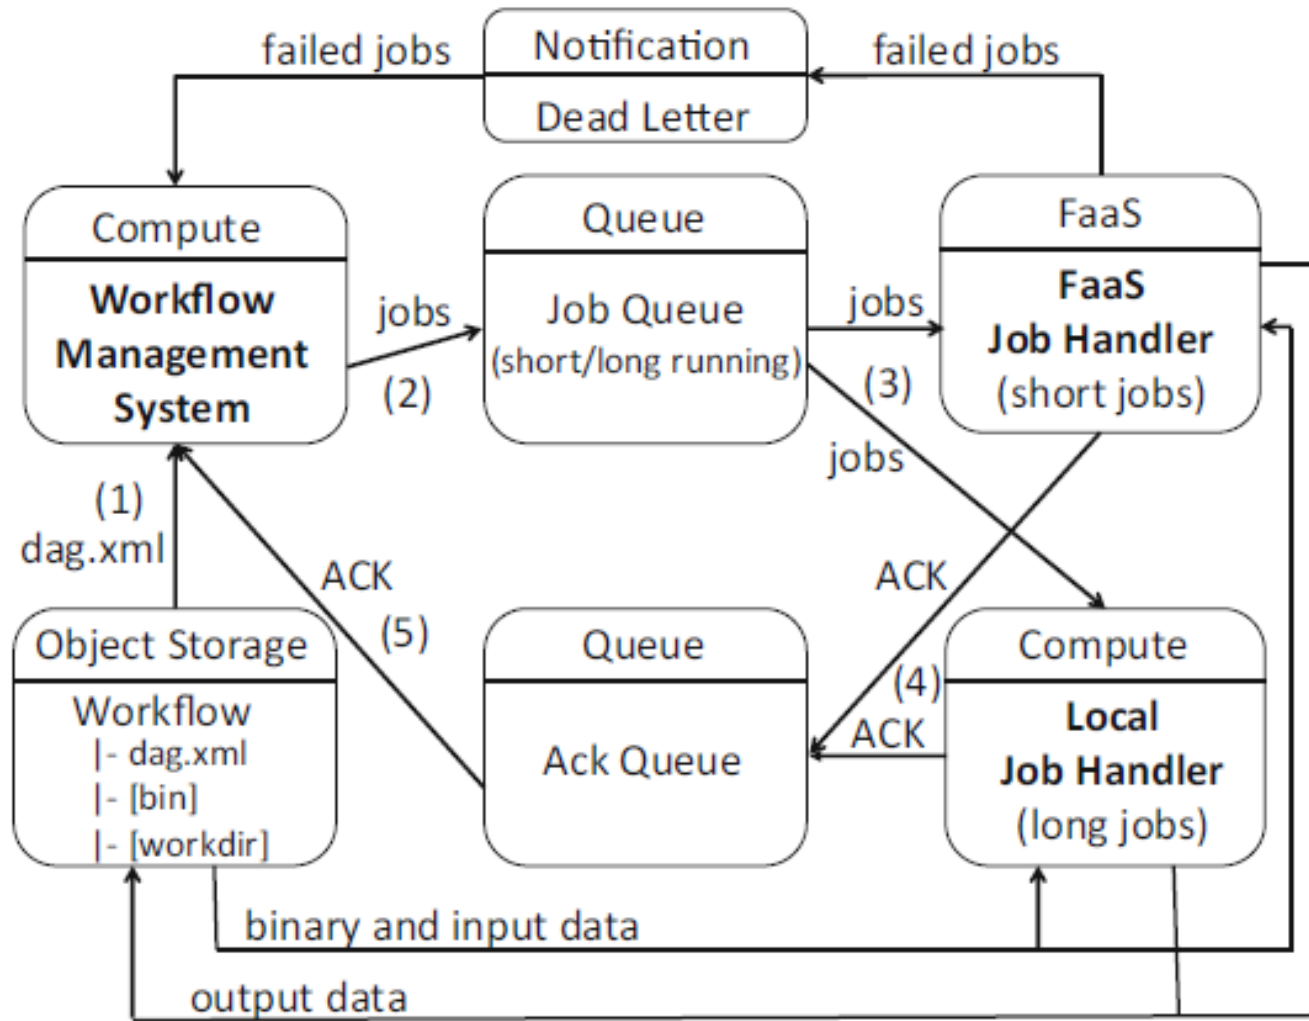
5 models to handle scientific workflows



- 📖 Jiang et al., created DEWE (Distributed Elastic Workflow Execution), a scheduler for running scalable scientific workflows with high performance.
- 📖 It distinguishes the jobs to Short jobs and Long jobs.

Approaches > Scientific Workflow

Death letter contains diagnostics data



DEWE

- 📖 Kijak et al., in their paper “Challenges for scheduling scientific workflows on cloud functions” introduced a trade-off between **deadline** and **budget**.
- 📖 This made possible by **Hyperflow** engine.

Approaches > Scientific Workflow

Require: DAG, time (D_{user}) and budget (B_{user})

- 1: Sort all tasks based on their level
- 2: **if** $B_{user} < Cost_{low}(DAG)$ **then**
- 3: **return** no possible schedule
- 4: **else if** $B_{user} > Cost_{high}(DAG)$ **then**
- 5: **return** schedule map on the most expensive resource
- 6: **end if**
- 7: Compute the sub-deadline value for each task
- 8: **while** there is unscheduled task **do**
- 9: **for** $r \in \text{resources}$ **do**
- 10: Calculate quality measure $Q(t_{cur}, r)$
- 11: **end for**
- 12: $r_{selected} \Leftarrow r$ with highest quality measure
- 13: assign t_{cur} to $r_{selected}$
- 14: **end while**
- 15: **return** schedule map

$$l(t_i) = 1 + \max_{t_p \in \text{predecessors}(t_i)} l(t_p),$$

$$Level_{execution}^j = \max_{l(t_i)=j} \{ET_{max}(t_i)\}$$

$$Level_{DL}^j = Level_{DL}^{j-1} + D_{user} * \frac{Level_{execution}^j}{\sum_{1 \leq j' \leq l(t_{exit})} Level_{execution}^{j'}}$$

$$S_{DL}(t_{cur}) = \{Level_{DL}^j | l(t_i) == j\}$$

$$Time_Q(t_{cur}, r) = \frac{\xi * S_{DL}(t_{cur}) - FT(t_{cur}, r)}{FT_{max}(t_{cur}) - FT_{min}(t_{cur})}$$

$$Cost_Q(t_{cur}, r) = \frac{Cost_{max}(t_{cur}) - Cost(t_{cur}, r)}{Cost_{max}(t_{cur}) - Cost_{min}(t_{cur})} * \xi$$

where

$$\xi = \begin{cases} 1 & \text{if } FT(t_{cur}, r) < S_{DL}(t_{cur}) \\ 0 & \text{otherwise} \end{cases}$$

$$Q(t_{cur}, r) = Time_Q(t_{cur}, r) * (1 - C_F) + Cost_Q(t_{cur}, r) * C_F$$

where C_F , a cost-efficient factor is a tradeoff factor defined as:

$$C_F = \frac{Cost_{low}(DAG)}{B_{user}}$$

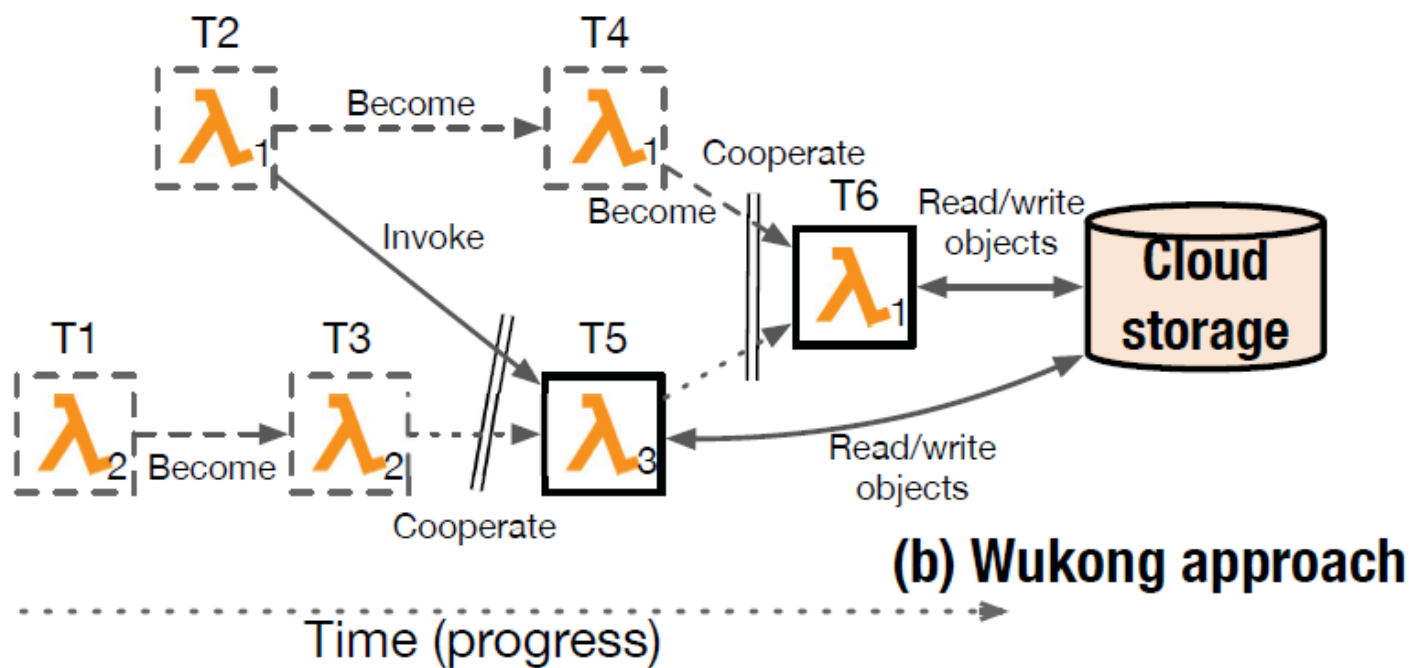
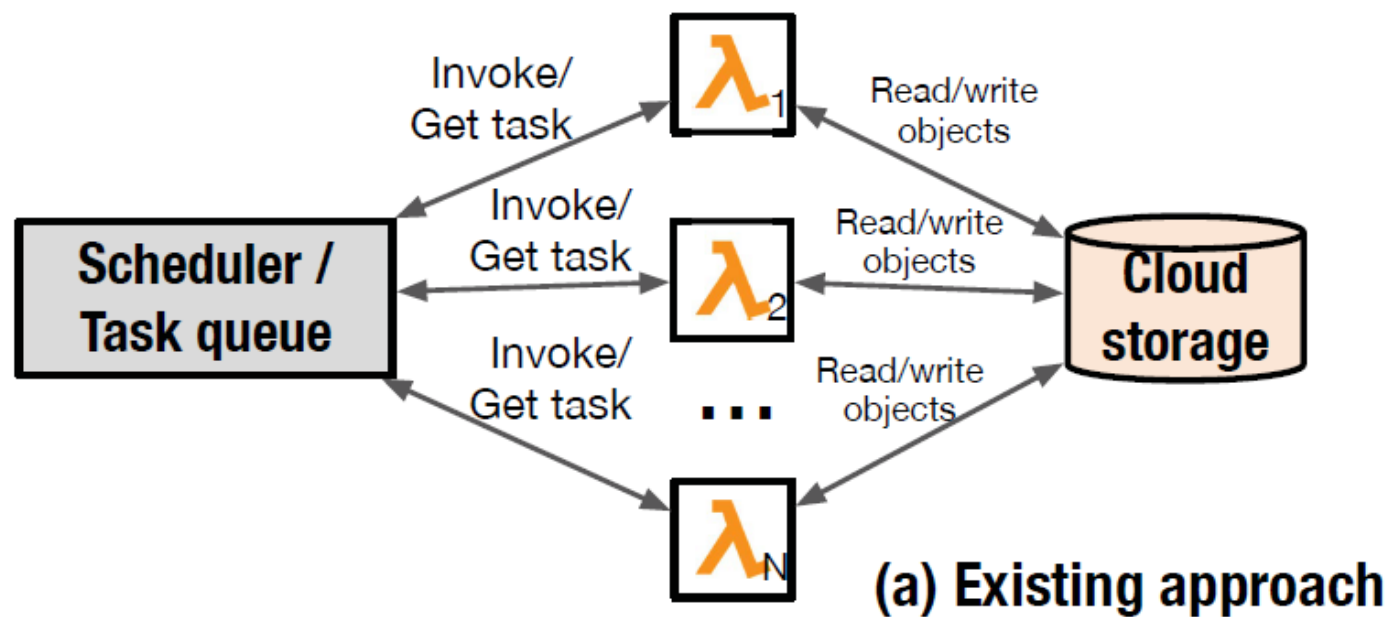
Parallel approach

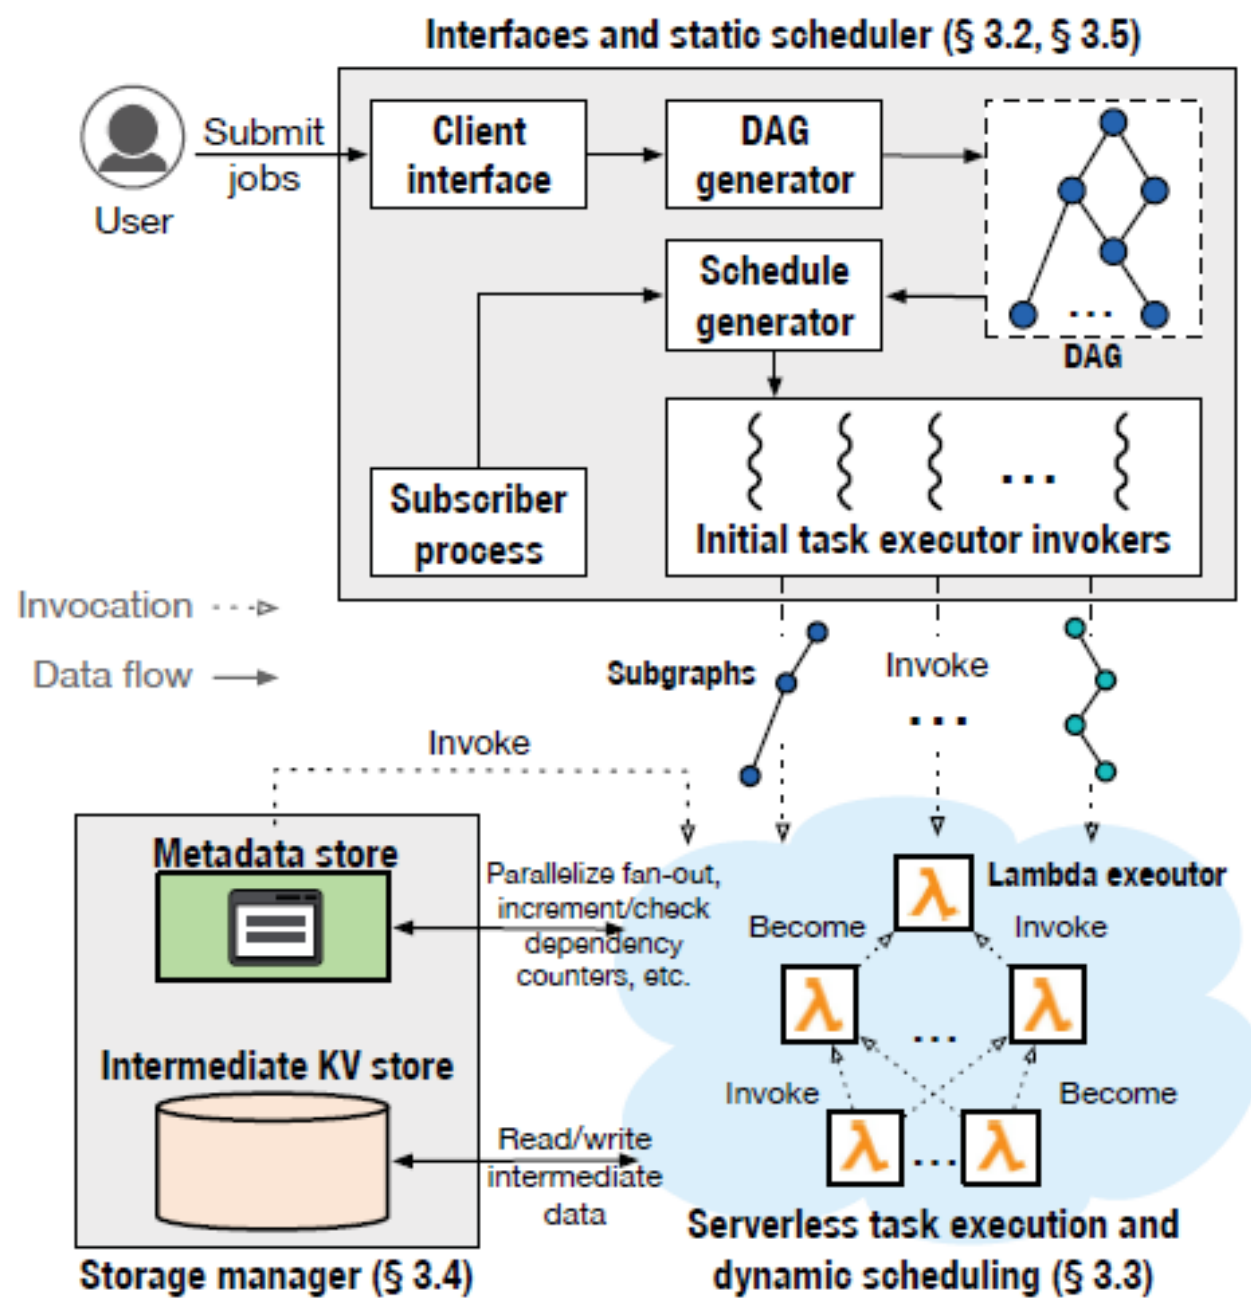
Approaches > Parallel


- Carver et al., introduced "Wukong", a **decentralized** approach to schedule serverless functions.
- Distributed scheduling enables us to run scheduling on all **runners in parallel**.
- This has 4 benefits:
 - 1. Data Locality usage optimization
 - 2. Less network I/O
 - 3. Auto scaling
 - 4. Cost efficiency

Approaches > Parallel

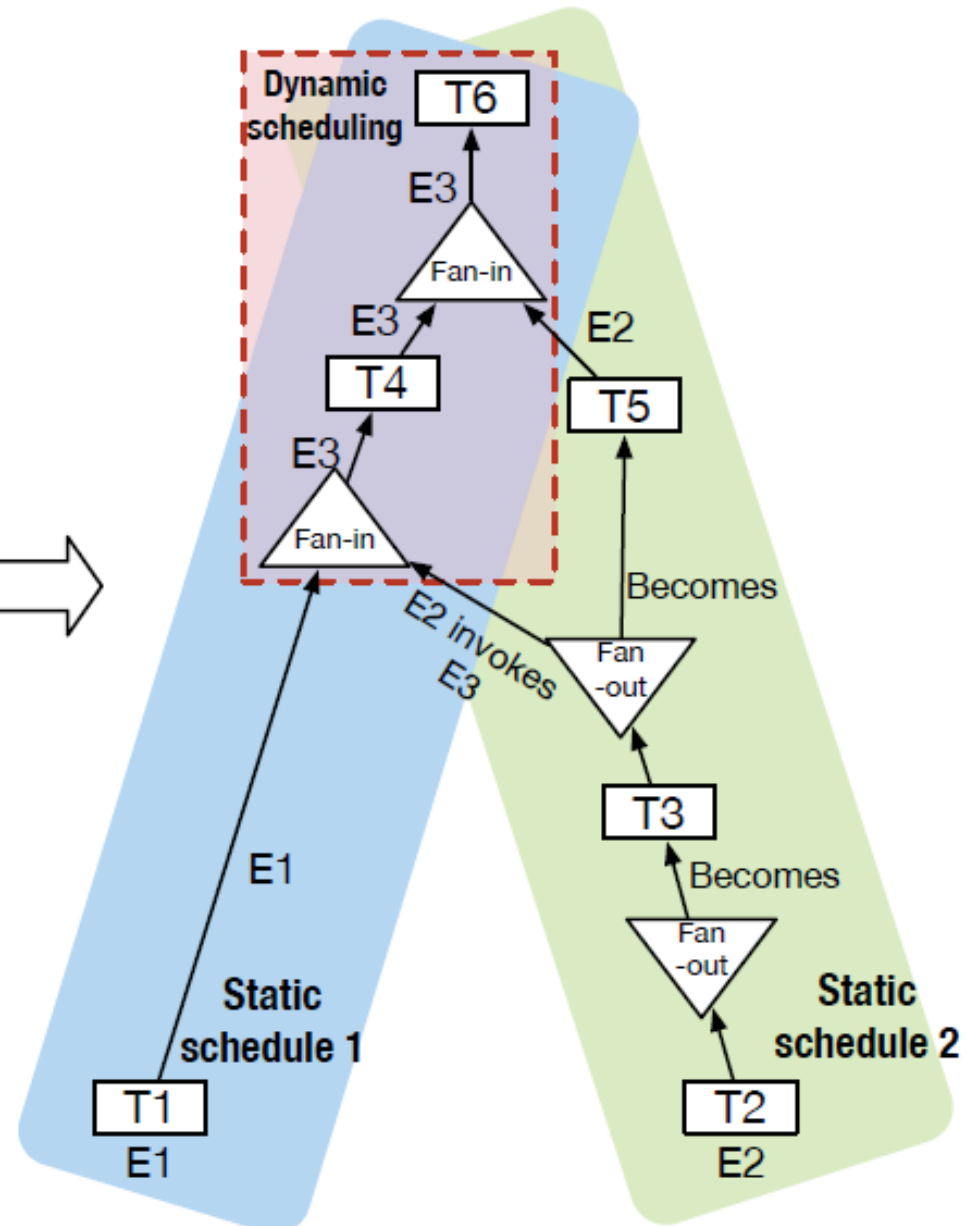
- 📖 They succeeded to run parallel tasks 68.17% faster
- 📖 They also reduced costs up to 92.96%







(a) Static DAG



(b) Dynamic scheduling

References

T. Harter et al., "Slacker: Fast distribution with lazy docker containers," in 14th USENIX Conference on File and Storage Technologies (FAST 16), 2016, Santa Clara, United States, pp. 181-195.


I.E. Akkus et al., "SAND: Towards High-Performance Serverless Computing," in Usenix Annual Technical Conference (18), July-2019, Boston, USA, pp. 923-935.

G. Aumala et al., "Beyond Load Balancing: Package-Aware Scheduling for Serverless Platforms," in 19th IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGRID) , May-2019, Larnaca, Cyprus, pp. 282-291.

M. Malawski et al., "Serverless execution of scientific workflows: Experiments with HyperFlow, AWS Lambda and Google Cloud Functions," Future Generation Computer Systems, vol. 110, pp. 502-514, September 2020.

J. Kijak et al., "Challenges for scheduling scientific workflows on cloud functions," in 2018 IEEE 11th International Conference on Cloud Computing (CLOUD), Jul-2018, San Francisco, USA, pp. 460-467.

B. Carver et al., "Wukong: a scalable and locality-enhanced framework for serverless parallel computing," in Proceedings of the 11th ACM Symposium on Cloud Computing, 2020, New York, USA, pp. 1-15.



Thank you!

Any questions?

☆ download here: <https://avestura.dev/presentations>

